

## QoS (Quality of Service)

### QoS Classification

QoS Classification is used to classify the packets to assign a Traffic Class (TC) value based on service requirements. TC is the internal priority in switch (8 possible value) using for packet buffering and scheduling. TC-to-priority-group (PG) and TC-to-Queue maps depend on the TC maps to the ingress buffer (priority-group) and egress queue, respectively. TC is used as a priority field for packet re-marking on egress port. Determines the TC based on the priority of the incoming packet and trust mode of a port.

Dot1p-to-TC Mapping is assigned the TC based on PCP (Priority Code Point) in the VLAN tag, and this mapping table is configured per ingress physical port.

Default Dot1p-to-TC mapping on Broadcom platform:

Dot1P	0	1	2	3	4	5	6	7
TC	0	1	2	3	4	5	6	7

Default Dot1p-to-TC mapping on the Intel platform:

Dot1P 0~7

TC 0

DSCP-to-TC Mapping is assigned the TC based on DSCP field in the IP header, and this mapping table is configured per ingress physical port.

Default DSCP-to-TC mapping:

DSCP 0~63

TC 0

TC-to-Queue Mapping is assigned the Queue ID based on TC, and this mapping table is configured per ingress physical port.

Default TC-to-Queue mapping on Broadcom platform:

TC	0	1	2	3	4	5	6	7
Queue	0	1	2	3	4	5	6	7

Default TC-to-Queue mapping on Intel platform:

TC 0~7

Queue 0

Tested model & firmware version:

Switch model name:

DCS503 (AS9716-32D)

Edgecore SONiC version:

202111.1

202111.3

202111.8

Topology:

mceclip1.png

Pre-configuration:

Add VLAN 10 to Ethernet0, Ethernet8.

admin@sonic:~\$ show vlan brief

```
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
| VLAN ID | IP Address | Ports | Port Tagging | Proxy ARP | DHCP
Helper | DHCP Source | DHCP Link | | | |
| Interface | Selection | | | | | Address
+=====+=====+=====+=====+=====+=====+
=====+=====+=====+
|      10 | | Ethernet0 | tagged | disabled | |
| | | | | | |
| | | Ethernet8 | tagged | | |
| | | | | |
```

```
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
```

Procedure:

Step 1: Create a profile for DOT1P/DSCP to TC(Traffic Class).

Example for Dot1P:

```
admin@sonic:~$ sudo config qos dot1p-tc add 1p_tc --dot1p 0 --tc 1
```

Example for DCSP:

```
admin@sonic:~$ sudo config qos dscp-tc add DSCP_TC --dscp 7 --tc 1
```

Note:

```
admin@sonic:~$ sudo config qos dot1p-tc add --help
```

Usage: config qos dot1p-tc add [OPTIONS] <profile>

Add dot1p-tc map profile.

Options:

```
--dot1p TEXT          Cos value [required]
--tc INTEGER RANGE    Traffic-class(TC) value [required]
-h, -?, --help        Show this message and exit.
```

Step 2: Modify the existing Dot1p/DSCP to TC profile.

Example for DOT1P:

```
admin@sonic:~$ sudo config qos dot1p-tc update 1p_tc --dot1p 1 --tc 2
```

Example for DCSP:

```
admin@sonic:~$ sudo config qos dscp-tc update DSCP_TC --dscp 8 --tc 2
```

Note:

```
admin@sonic:~$ sudo config qos dot1p-tc update --help
```

Usage: config qos dot1p-tc update [OPTIONS] <profile>

Update dot1p-tc map profile.

Options:

```
--dot1p TEXT          Cos value [required]
--tc INTEGER RANGE    Traffic-class(TC) value
--remove              Delete the mapping entry
-?, -h, --help        Show this message and exit.
```

Step 3: Check the profile of DOT1P/DSCP to TC.

DOT1P to TC:

```
admin@sonic:~$ show qos dot1p-tc
```

dot1p-tc policy: 1p\_tc

```
  Dot1p    TC
  -----
      0      1
      1      2
```

DSCP to TC:

```
admin@sonic:~$ show qos dscp-tc
```

dscp-tc policy: DSCP\_TC

```
  DSCP    TC
  -----
      7      1
      8      2
```

Step 4: Create a profile for TC to Queue.

```
admin@sonic:~$ sudo config qos tc-queue add TC_Q --tc 1 --queue 2
Note:
```

```
admin@sonic:~$ sudo config qos tc-queue add --help
Usage: config qos tc-queue add [OPTIONS] <profile>
```

Add tc-queue map profile.

Options:

```
--tc TEXT           Traffic-class(TC) value [required]
--queue INTEGER RANGE Queue value [required]
-?, -h, --help      Show this message and exit.
```

Step 5: Modify the existing TC to Queue profile.

```
admin@sonic:~$ sudo config qos tc-queue update TC_Q --tc 2 --queue 3
Note:
```

```
admin@sonic:~$ sudo config qos tc-queue update --help
Usage: config qos tc-queue update [OPTIONS] <profile>
```

Update tc-queue map profile.

Options:

```
--tc TEXT           Traffic-class(TC) value [required]
--queue INTEGER RANGE Queue value
--remove            Delete the mapping entry
-?, -h, --help      Show this message and exit.
```

Step 6: Check the profile of TC to Queue.

```
admin@sonic:~$ show qos tc-queue
tc-queue policy: TC_Q
  TC      Queue
  ----  -
    1      2
    2      3
```

Step 7: Binding the mapping table to the specified interface.

```
admin@sonic:~$ sudo config interface qos dot1p-tc bind Ethernet0 1p_tc
admin@sonic:~$ sudo config interface qos dscp-tc bind Ethernet0 DSCP_TC
admin@sonic:~$ sudo config interface qos tc-queue bind Ethernet0 TC_Q
Note:
```

```
admin@sonic:~$ sudo config interface qos dot1p-tc --help
Usage: config interface qos dot1p-tc [OPTIONS] <op> <interface_name> <profile>
```

dot1p-tc policy configuration

Options:

```
-?, -h, --help  Show this message and exit.
"op": "bind/unbind"
```

Step 6: Check the binding table.

```
admin@sonic:~$ show interfaces qos
Ethernet0:
```

```

Dot1p to TC: 1p_tc
DSCP to TC: DSCP_TC
TC to Queue: TC_Q

```

Result: Check the queue counters

When the VLAN 10 packet with priority 0 into the switch, it will follow the mapping table that we set.

Dot1p ----- > TC ----- > Queue

Based on the above setting, the VLAN 10 packet with priority 0 will be forwarded to Queue2.

Step 1. Clear the queue counter

```
admin@sonic:~$ sonic-clear queuecounters
```

Step 2. Check Ethernet8 (egress port) queue counter

```
admin@sonic:~$ show queue counters Ethernet8
```

Last cached time was 2022-09-13 21:39:53.822559

Port	TxQ	Counter/pkts	Counter/bytes	Drop/pkts	Drop/bytes
Ethernet8	UC0	0	0	0	0
Ethernet8	UC1	0	0	0	0
Ethernet8	UC2	37,846	2,422,144	0	0
Ethernet8	UC3	0	0	0	0
Ethernet8	UC4	0	0	0	0
Ethernet8	UC5	0	0	0	0
Ethernet8	UC6	0	0	0	0
Ethernet8	UC7	0	0	0	0
Ethernet8	UC8	0	0	0	0
Ethernet8	UC9	0	0	0	0
Ethernet8	MC10	0	0	0	0
Ethernet8	MC11	0	0	0	0

QoS Marking

Once the packets have been classified by QoS Classification, QoS Marking allows to set or change the priority value in the packets before sending them. The new priority value is brought to the next switches for fine-tuning or uniform the class of the traffic in the network.

QoS Marking enables to rewrite the Dot1p or DSCP of incoming packets based on Traffic Class (TC) value.

Enable Dot1p rewrite by configuring a TC-to-Dot1p mapping on an egress port. On Broadcom switches, the Dot1p rewrite is always enabled.

Default TC-to-Dot1p mapping.

TC	0	1	2	3	4	5	6	7
Dot1P	0	1	2	3	4	5	6	7

Enable DSCP rewrite by configuring a TC-to-DSCP mapping on an egress port. The DSCP rewrite is disabled by default on all ports.

Tested model & firmware version:

Switch model name:

DCS503 (AS9716-32D)

Edgecore SONiC version:

202111.1

202111.3

Topology:

mceclip0.png

Pre-configuration:  
Add VLAN 10 to Ethernet0, Ethernet8.

admin@sonic:~\$ show vlan brief

+-----+-----+-----+-----+-----+-----+-----+						
+-----+-----+-----+-----+-----+-----+-----+						
VLAN ID	IP Address	Ports	Port Tagging	Proxy ARP	DHCP	
Helper	DHCP Source	DHCP Link				
Interface	Selection				Address	
+=====+=====+=====+=====+=====+=====+=====+						
10		Ethernet0	tagged	disabled		
		Ethernet8	tagged			
+-----+-----+-----+-----+-----+-----+-----+						
+-----+-----+-----+-----+-----+-----+-----+						

Procedure:  
Step 1: Create a profile for DOT1P remarking.

Example for Dot1P:

admin@sonic:~\$ sudo config qos remark dot1p add remark\_dot1p --tc 0 --dot1p 1  
Example for DSCP:

admin@sonic:~\$ sudo config qos remark dscp add remark\_dscp --tc 0 --dscp 7  
Note:

admin@sonic:~\$ sudo config qos remark dot1p add --help  
Usage: config qos remark dot1p add [OPTIONS] <profile>

Add tc-dot1p map profile.

Options:

--tc TEXT Traffic-class(TC) value [required]  
--dot1p INTEGER RANGE Cos value [required]  
-?, -h, --help Show this message and exit.

Step 2: Modify the existing remark profile.

Example for Dot1P:

admin@sonic:~\$ sudo config qos remark dot1p update remark\_dot1p --tc 1 --dot1p 2  
Example for DSCP:

admin@sonic:~\$ sudo config qos remark dscp update remark\_dscp --tc 1 --dscp 8  
Note:

admin@sonic:~\$ sudo config qos remark dot1p update --help  
Usage: config qos remark dot1p update [OPTIONS] <profile>

Update tc-dot1p map profile.

Options:

--tc TEXT Traffic-class(TC) value [required]  
--dot1p INTEGER RANGE Cos value  
--remove Delete the mapping entry  
-h, -?, --help Show this message and exit.

Step 3: Check the remark profile.

admin@sonic:~\$ show qos remark dot1p  
dot1p policy: remark\_dot1p  
TC Dot1p

```

-----
0      1
1      2

```

Step 4: Binding the remark table to the egress interface.

```

admin@sonic:~$ sudo config interface qos remark dot1p bind Ethernet8
remark_dot1p

```

Note:

```

admin@sonic:~$ sudo config interface qos remark dot1p --help
Usage: config interface qos remark dot1p [OPTIONS] <op> <interface_name>
      <profile>

```

tc-dot1p policy configuration

Options:

-?, -h, --help Show this message and exit.

"op": "bind/unbind"

Step 5: Check the binding table.

```

admin@sonic:~$ show interfaces qos

```

Ethernet8:

Dot1p remark: remark\_dot1p

Result:

Since this example didn't set the QoS Classification, when the VLAN 10 packet with priority 0 into switch, it will base on the default QoS Classification. The priority 0 will map to TC 0.

Base on the remark profile, the TC 0 will map to the Dot1p 1 which means the traffic through the TC 0 will be remarked to priority 1.

Here's the packet.

mceclip2.png

## QoS Scheduler

Queue scheduling provides preferential treatment of traffic classes mapped to specific egress queues. SONiC supports SP, WRR, and DWRR scheduling disciplines.

SP (Strict Priority)- Higher priority egress queues get scheduled for transmission over lower priority queues.

WRR (Weighted Round Robin) - Egress queues receive bandwidth proportional to the configured weight. The scheduling granularity is per packet which causes large and small packets to be treated the same. Flows with large packets have an advantage over flows with smaller packets.

DWRR (Deficit Weighted Round Robin) - Similar to WRR but uses deficit counter scheduling granularity to account for packet size variations and provide a more accurate proportion of bandwidth.

Tested model & firmware version:

Switch model name:

DCS203 (AS7326-56X)

Edgecore SONiC version:

202111.1

202111.3

Topology:

mceclip3.png

Pre-configuration:

Add VLAN 10 to Ethernet0, Ethernet1, and Ethernet5.

```

admin@sonic:~$ show vlan brief

```

```

+-----+-----+-----+-----+-----+

```

VLAN ID		IP Address	Ports	Port Tagging	Proxy ARP	DHCP
Helper	DHCP Source	DHCP Link				Address
Interface	Selection					
+-----+-----+-----+-----+-----+-----+-----						
10		Ethernet0	tagged	disabled		
		Ethernet1	tagged			
		Ethernet5	tagged			
+-----+-----+-----+-----+-----+-----+-----						

The QoS Classification uses the default setting.

SP (Strict Priority)

Procedure:

Step 1: Set the scheduler mode.

```
admin@sonic:~$ sudo config scheduler add strict_mode --sched_type STRICT
```

Note:

```
admin@sonic:~$ sudo config scheduler add --help
```

```
Usage: config scheduler add [OPTIONS] <profile_name>
```

Add QoS-Scheduler profile.

Options:

```
--sched_type [WRR|DWRR|STRICT] Scheduler type. WRR is not supported on
                                Intel platform.
--weight INTEGER RANGE          weight
--shaper_type [bytes|packets]    Shaper type
--bandwidth TEXT                 Maximum bandwidth rate in kbps or pps. If
                                shaper-type is bytes, user can specify a
                                decimal number followed by the abbreviation
                                k (1000), m (1,000,000), or g
                                (1,000,000,000)
-h, -?, --help                  Show this message and exit.
```

Step 2: Check status.

```
admin@sonic:~$ show scheduler
```

Name	Scheduling Type	Weight	Shaper Type	Bandwidth
strict_mode	STRICT	N/A	N/A	N/A

Step 3: Binding the scheduler to the interface.

```
admin@sonic:~$ sudo config interface scheduler bind queue Ethernet5 3
```

```
strict_mode
```

Note:

```
admin@sonic:~$ sudo config interface scheduler bind queue --help
```

```
Usage: config interface scheduler bind queue [OPTIONS] <interface_name>
                                <queue> <sched_policy_name>
```

Options:

```
-h, -?, --help Show this message and exit.
```

Here's the command to unbind the scheduler.

```
admin@sonic:~$ sudo config interface scheduler unbind queue Ethernet5 3
```

Step 4: Check status.

```
admin@sonic:~$ show interfaces scheduler
```

```
Ethernet5
```

Type	Queue	Profile	Scheduling Mode	Weight	Shaper Type
------	-------	---------	-----------------	--------	-------------

Bandwidth					
-----------	--	--	--	--	--

Queue	3	strict_mode	STRICT	N/A	N/A
-------	---	-------------	--------	-----	-----

```
0  
Result:
```

Injecting the traffic with priority 0~7 from Ethernet0 and Ethernet1 to Ethernet5 to trigger the congestion.

Queue 3 in Ethernet5 will be the highest priority.

```
admin@sonic:~$ show queue counters Ethernet5
```

```
Last cached time was 2022-09-14 06:40:04.548068
```

Port	TxQ	Counter/pkts	Counter/bytes	Drop/pkts	Drop/bytes
Ethernet5	UC0	1,824,146	233,490,688	2,407,299	308,134,272
Ethernet5	UC1	1,824,176	233,494,528	2,407,270	308,130,560
Ethernet5	UC2	1,824,179	233,494,912	2,407,292	308,133,376
Ethernet5	UC3	4,255,417	544,693,376	0	0
Ethernet5	UC4	1,824,185	233,495,680	2,407,305	308,135,040
Ethernet5	UC5	1,824,187	233,495,936	2,407,276	308,131,328
Ethernet5	UC6	1,824,189	233,496,192	2,407,272	308,130,816
Ethernet5	UC7	1,824,194	233,496,832	2,407,284	308,132,352
Ethernet5	UC8	0	0	0	0
Ethernet5	UC9	0	0	0	0
Ethernet5	MC10	0	0	0	0
Ethernet5	MC11	0	0	0	0
Ethernet5	MC12	0	0	0	0
Ethernet5	MC13	0	0	0	0
Ethernet5	MC14	0	0	0	0
Ethernet5	MC15	0	0	0	0
Ethernet5	MC16	0	0	0	0
Ethernet5	MC17	0	0	0	0
Ethernet5	MC18	0	0	0	0
Ethernet5	MC19	0	0	0	0

```
Notes:
```

8 dedicated queues (i.e UC0, UC1, ..., UC7) for unicast.

Other 8 dedicated queues (i.e MC10, MC11, ..., MC17) for multicast.

UC8, UC9, UC18, UC19 are not available.

WRR (Weighted Round Robin)

Procedure:

Step 1: Set the scheduler mode.

```
admin@sonic:~$ sudo config scheduler add wrr_7 --sched_type WRR --weight 7
```

```
admin@sonic:~$ sudo config scheduler add wrr_3 --sched_type WRR --weight 3
```

Step 2: Check status.

```
admin@sonic:~$ show scheduler
```

Name	Scheduling Type	Weight	Shaper Type	Bandwidth
wrr_3	WRR	3	N/A	N/A
wrr_7	WRR	7	N/A	N/A

Step 3: Binding the scheduler to the interface.

```
admin@sonic:~$ sudo config interface scheduler bind queue Ethernet5 3 wrr_7
```

```
admin@sonic:~$ sudo config interface scheduler bind queue Ethernet5 4 wrr_3
```

Step 4: Check status.

```
admin@sonic:~$ show interfaces scheduler
```



Ethernet5 Type Bandwidth	Queue	Profile	Scheduling Mode	Weight	Shaper Type
--------------------------------	-------	---------	-----------------	--------	-------------

Queue 0	3	wrr_7	WRR	7	N/A
Queue 0	4	wrr_3	WRR	3	N/A

Result:

Injecting the traffic with priority 3,4 from Ethernet0 and Ethernet1 to Ethernet5 to trigger the congestion.  
The egress traffic will be distributed based on weights.

admin@sonic:~\$ show queue counters Ethernet5

Last cached time was 2022-09-23 05:53:03.667274

Port	TxQ	Counter/pkts	Counter/bytes	Drop/pkts	Drop/bytes
Ethernet5	UC0	1	409	0	0
Ethernet5	UC1	0	0	0	0
Ethernet5	UC2	0	0	0	0
Ethernet5	UC3	1,518,614	858,150,900	620,044	348,958,020
Ethernet5	UC4	671,393	379,639,721	1,467,263	827,705,931
Ethernet5	UC5	0	0	0	0
Ethernet5	UC6	0	0	0	0
Ethernet5	UC7	0	0	0	0
Ethernet5	UC8	0	0	0	0
Ethernet5	UC9	0	0	0	0
Ethernet5	MC10	0	0	0	0
Ethernet5	MC11	0	0	0	0
Ethernet5	MC12	0	0	0	0
Ethernet5	MC13	0	0	0	0
Ethernet5	MC14	0	0	0	0
Ethernet5	MC15	0	0	0	0
Ethernet5	MC16	0	0	0	0
Ethernet5	MC17	0	0	0	0
Ethernet5	MC18	0	0	0	0
Ethernet5	MC19	0	0	0	0

Note:

Total: 1,518,614 + 671,393 = 2,190,007

Queue 3: 1,518,614 / 2,190,007 = 0.6934

Queue 4: 671,393 / 2,190,007 = 0.3065

DWRR (Deficit Weighted Round Robin)

Procedure:

Step 1: Set the scheduler mode.

admin@sonic:~\$ sudo config scheduler add dwrr\_7 --sched\_type DWRR --weight 7

admin@sonic:~\$ sudo config scheduler add dwrr\_3 --sched\_type DWRR --weight 3

Step 2: Check status.

admin@sonic:~\$ show scheduler

Name	Scheduling Type	Weight	Shaper Type	Bandwidth
dwrr_3	DWRR	3	N/A	N/A
dwrr_7	DWRR	7	N/A	N/A

Step 3: Binding the scheduler to the interface.

admin@sonic:~\$ sudo config interface scheduler bind queue Ethernet5 3 dwrr\_7

admin@sonic:~\$ sudo config interface scheduler bind queue Ethernet5 4 dwrr\_3

Step 4: Check status.

```
admin@sonic:~$ show interfaces scheduler
```

Ethernet5

Type	Queue	Profile	Scheduling Mode	Weight	Shaper Type
------	-------	---------	-----------------	--------	-------------

Bandwidth					
-----------	--	--	--	--	--

Queue 0	3	dwrr_7	DWRR	7	N/A
---------	---	--------	------	---	-----

Queue 0	4	dwrr_3	DWRR	3	N/A
---------	---	--------	------	---	-----

0

Result:

Injecting the traffic with priority 3,4 from Ethernet0 and Ethernet1 to Ethernet5 to trigger the congestion.

The egress traffic will be distributed based on weights.

```
admin@sonic:~$ show queue counters Ethernet5
```

Last cached time was 2022-09-28 07:07:32.975167

Port	TxQ	Counter/pkts	Counter/bytes	Drop/pkts	Drop/bytes
------	-----	--------------	---------------	-----------	------------

Ethernet5	UC0	3	1,227	0	0
-----------	-----	---	-------	---	---

Ethernet5	UC1	0	0	0	0
-----------	-----	---	---	---	---

Ethernet5	UC2	0	0	0	0
-----------	-----	---	---	---	---

Ethernet5	UC3	5,972,663	764,500,864	2,473,283	316,580,224
-----------	-----	-----------	-------------	-----------	-------------

Ethernet5	UC4	2,605,564	333,512,192	5,840,382	747,568,896
-----------	-----	-----------	-------------	-----------	-------------

Ethernet5	UC5	0	0	0	0
-----------	-----	---	---	---	---

Ethernet5	UC6	0	0	0	0
-----------	-----	---	---	---	---

Ethernet5	UC7	0	0	0	0
-----------	-----	---	---	---	---

Ethernet5	UC8	0	0	0	0
-----------	-----	---	---	---	---

Ethernet5	UC9	0	0	0	0
-----------	-----	---	---	---	---

Ethernet5	MC10	0	0	0	0
-----------	------	---	---	---	---

Ethernet5	MC11	0	0	0	0
-----------	------	---	---	---	---

Ethernet5	MC12	0	0	0	0
-----------	------	---	---	---	---

Ethernet5	MC13	0	0	0	0
-----------	------	---	---	---	---

Ethernet5	MC14	0	0	0	0
-----------	------	---	---	---	---

Ethernet5	MC15	0	0	0	0
-----------	------	---	---	---	---

Ethernet5	MC16	0	0	0	0
-----------	------	---	---	---	---

Ethernet5	MC17	0	0	0	0
-----------	------	---	---	---	---

Ethernet5	MC18	0	0	0	0
-----------	------	---	---	---	---

Ethernet5	MC19	0	0	0	0
-----------	------	---	---	---	---

Note:

Total: 5,972,663 + 2,605,564 = 8,578,227

Queue 3: 5,972,412 / 8,577,474 = 0.6962

Queue 4: 2,605,062 / 8,577,474 = 0.3037

Hybrid mode (SP + WRR)

Procedure:

Step 1: Set the scheduler mode.

```
admin@sonic:~$ sudo config scheduler add strict_mode --sched_type STRICT
```

```
admin@sonic:~$ sudo config scheduler add wr_7 --sched_type WRR --weight 7
```

```
admin@sonic:~$ sudo config scheduler add wr_3 --sched_type WRR --weight 3
```

Step 2: Check status.

```
admin@sonic:~$ show scheduler
```

Name	Scheduling Type	Weight	Shaper Type	Bandwidth
------	-----------------	--------	-------------	-----------

strict_mode	STRICT	N/A	N/A	N/A
-------------	--------	-----	-----	-----

wrr_3	WRR	3	N/A	N/A
-------	-----	---	-----	-----

wrr\_7                      WRR                                      7   N/A                                      N/A

Step 3: Binding the scheduler to the interface.

```
admin@sonic:~$ sudo config interface scheduler bind queue Ethernet5 2
strict_mode
admin@sonic:~$ sudo config interface scheduler bind queue Ethernet5 3 wrr_7
admin@sonic:~$ sudo config interface scheduler bind queue Ethernet5 4 wrr_3
Step 4: Check status.
```

```
admin@sonic:~$ show interfaces scheduler
Ethernet5
Type      Queue  Profile      Scheduling Mode  Weight  Shaper Type
Bandwidth
-----
Queue 2    strict_mode  STRICT          N/A      N/A
0
Queue 3    wrr_7       WRR             7        N/A
0
Queue 4    wrr_3       WRR             3        N/A
0
```

Result:

Injecting the traffic with priority 2,3,4 from Ethernet0 and Ethernet1 to Ethernet5 to trigger the congestion.  
Queue 2 in Ethernet5 will be the highest priority.  
Queue 3, and 4 in Ethernet 5 will distribute the remaining traffic based on weights.

```
admin@sonic:~$ show queue counters Ethernet5
Last cached time was 2022-09-23 05:39:48.815061
Port      TxQ      Counter/pkts  Counter/bytes  Drop/pkts  Drop/bytes
-----
Ethernet5 UC0        1           409            0            0
Ethernet5 UC1         0            0            0            0
Ethernet5 UC2    5,630,631   720,720,768    0            0
Ethernet5 UC3    2,035,749   260,575,872   3,594,882   460,144,896
Ethernet5 UC4    912,664    116,820,992   4,717,966   603,899,648
Ethernet5 UC5         0            0            0            0
Ethernet5 UC6         0            0            0            0
Ethernet5 UC7         0            0            0            0
Ethernet5 UC8         0            0            0            0
Ethernet5 UC9         0            0            0            0
Ethernet5 MC10        0            0            0            0
Ethernet5 MC11        0            0            0            0
Ethernet5 MC12        0            0            0            0
Ethernet5 MC13        0            0            0            0
Ethernet5 MC14        0            0            0            0
Ethernet5 MC15        0            0            0            0
Ethernet5 MC16        0            0            0            0
Ethernet5 MC17        0            0            0            0
Ethernet5 MC18        0            0            0            0
Ethernet5 MC19        0            0            0            0
```

Note:

Total: 2,035,749 + 912,664 = 2,948,413  
Queue 3: 2,035,749 / 2,948,413 = 0.6904  
Queue 4: 912,664 / 2,948,413 = 0.3095

Hybrid mode (SP + DWRR)

Procedure:

Step 1: Set the scheduler mode.

```

admin@sonic:~$ sudo config scheduler add strict_mode --sched_type STRICT
admin@sonic:~$ sudo config scheduler add dwrr_7 --sched_type DWRR --weight 7
admin@sonic:~$ sudo config scheduler add dwrr_3 --sched_type DWRR --weight 3
Step 2: Check status.

```

```

admin@sonic:~$ show scheduler
Name           Scheduling Type      Weight  Shaper Type      Bandwidth
-----
dwrr_3         DWRR                  3      N/A              N/A
dwrr_7         DWRR                  7      N/A              N/A
strict_mode    STRICT               N/A    N/A              N/A
Step 3: Binding the scheduler to the interface.

```

```

admin@sonic:~$ sudo config interface scheduler bind queue Ethernet5 2
strict_mode
admin@sonic:~$ sudo config interface scheduler bind queue Ethernet5 3 dwrr_7
admin@sonic:~$ sudo config interface scheduler bind queue Ethernet5 4 dwrr_3
Step 4: Check status.

```

```

admin@sonic:~$ show interfaces scheduler
Ethernet5
Type Queue Profile Scheduling Mode Weight Shaper Type
Bandwidth
-----
Queue 2 strict_mode STRICT N/A N/A
0
Queue 3 dwrr_7 DWRR 7 N/A
0
Queue 4 dwrr_3 DWRR 3 N/A
0
Result:

```

Injecting the traffic with priority 2,3,4 from Ethernet0 and Ethernet1 to Ethernet5 to trigger the congestion.  
Queue 2 in Ethernet5 will be the highest priority.  
Queue 3, and 4 in Ethernet 5 will distribute the remaining traffic based on weights.

```

admin@sonic:~$ show queue counters Ethernet5
Last cached time was 2022-09-29 02:36:49.112556
Port TxQ Counter/pkts Counter/bytes Drop/pkts Drop/bytes
-----
Ethernet5 UC0 8 2,989 0 0
Ethernet5 UC1 0 0 0 0
Ethernet5 UC2 1,898,110 242,958,080 0 0
Ethernet5 UC3 1,180,033 667,767,452 718,077 403,717,686
Ethernet5 UC4 521,926 295,427,164 1,376,184 776,243,178
Ethernet5 UC5 0 0 0 0
Ethernet5 UC6 0 0 0 0
Ethernet5 UC7 0 0 0 0
Ethernet5 UC8 0 0 0 0
Ethernet5 UC9 0 0 0 0
Ethernet5 MC10 0 0 0 0
Ethernet5 MC11 0 0 0 0
Ethernet5 MC12 0 0 0 0
Ethernet5 MC13 0 0 0 0
Ethernet5 MC14 0 0 0 0
Ethernet5 MC15 0 0 0 0
Ethernet5 MC16 0 0 0 0
Ethernet5 MC17 0 0 0 0
Ethernet5 MC18 0 0 0 0
Ethernet5 MC19 0 0 0 0

```

Note:

Total:  $1,180,033 + 521,926 = 1,701,959$   
Queue 3:  $1,180,033 / 1,701,959 = 0.6933$   
Queue 4:  $521,926 / 1,701,959 = 0.3066$

### QoS Shaping

Queue shaping provides control of minimum and maximum bandwidth requirements per egress queue for more effective bandwidth utilization. (Note: There is no minimum bandwidth guarantee support for queue shaping on Broadcom and Intel switches.) Egress queues that exceed an average transmission rate beyond the shaper max bandwidth will stop being serviced. Additional ingress traffic will continue to be stored on the egress queue until the queue size is exceeded which results in tail drop.

Tested model & firmware version:

Switch model name:

DCS203 (AS7326-56X)

Edgecore SONiC version:

202111.1

202111.3

Topology:

mceclip4.png

Pre-configuration:

Add VLAN 10 to Ethernet0, Ethernet1, Ethernet5.

```
admin@sonic:~$ show vlan brief
```

+-----+-----+-----+-----+-----+-----+-----						
+-----+-----+-----+-----+-----+-----+-----						
VLAN ID	IP Address	Ports	Port Tagging	Proxy ARP	DHCP	
Helper	DHCP Source	DHCP Link				
Interface	Selection				Address	
+=====+=====+=====+=====+=====+=====+=====						
10		Ethernet0	tagged	disabled		
		Ethernet1	tagged			
		Ethernet5	tagged			
+-----+-----+-----+-----+-----+-----+-----						
+-----+-----+-----+-----+-----+-----+-----						

Procedure:

Step 1: Create the profile for QoS shaping.

```
admin@sonic:~$ sudo config scheduler add 400m_shaping --shaper_type bytes --bandwidth 400m
```

Note:

```
admin@sonic:~$ sudo config scheduler add --help
```

Usage: config scheduler add [OPTIONS] <profile\_name>

Add QoS-Scheduler profile.

Options:

--sched\_type [WRR|DWRR|STRICT] Scheduler type. WRR is not supported on Intel platform.  
--weight INTEGER RANGE weight

--shaper\_type [bytes|packets] Shaper type  
--bandwidth TEXT Maximum bandwidth rate in kbps or pps. If shaper-type is bytes, user can specify a decimal number followed by the abbreviation k (1000), m (1,000,000), or g (1,000,000,000)  
-h, -?, --help Show this message and exit.

Step 2: Check the profile:

```
admin@sonic:~$ show scheduler
```

Name	Scheduling Type	Weight	Shaper Type	Bandwidth
400m_shaping	N/A	N/A	bytes	400 Mbps

Step 3: Binding the QoS shaping to the egress interface.

```
admin@sonic:~$ sudo config interface scheduler bind queue Ethernet5 3
```

```
400m_shaping
```

Note:

Usage: config interface scheduler <bind/unbind> queue [OPTIONS] <interface\_name> <queue> <sched\_policy\_name>

Options:

-?, -h, --help Show this message and exit.

Or

Usage: config interface scheduler <bind/unbind> port [OPTIONS] <interface\_name> <sched\_policy\_name>

Options:

-h, -?, --help Show this message and exit.

Step 4: Check the binding table.

```
admin@sonic:~$ show interfaces scheduler
```

Ethernet5	Type	Queue	Profile	Scheduling Mode	Weight	Shaper Type	Bandwidth
		3	400m_shaping	N/A	N/A	bytes	400 Mbps

Result:

Since the QoS shaping, the traffic through to the Queue 3 will be down to about 400 Mbps.

mceclip5.png